

Machine Learning and Deep Learning

Lecture-05

By: Somnath Mazumdar
Assistant Professor
sma.digi@cbs.dk

Overview

- Dimensionality Reduction
- Principal Component Analysis (PCA)
- Singular value decomposition (SVD) [Not in Lecture Plan]
- Decision Trees
- Random Forests

Dimensionality Reduction

Dimension

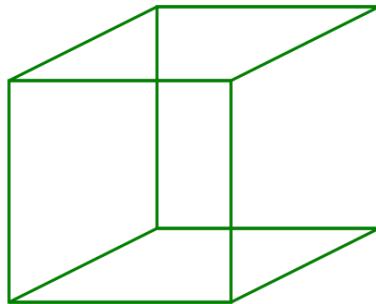
1 Dimension



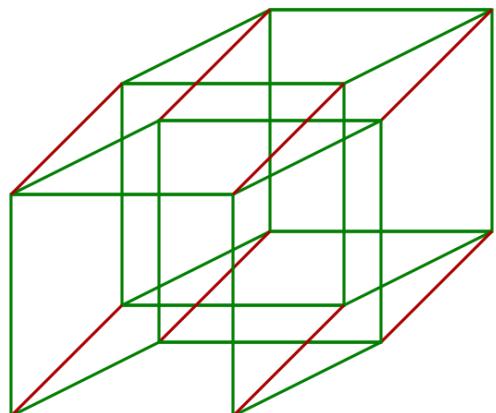
2 Dimensions



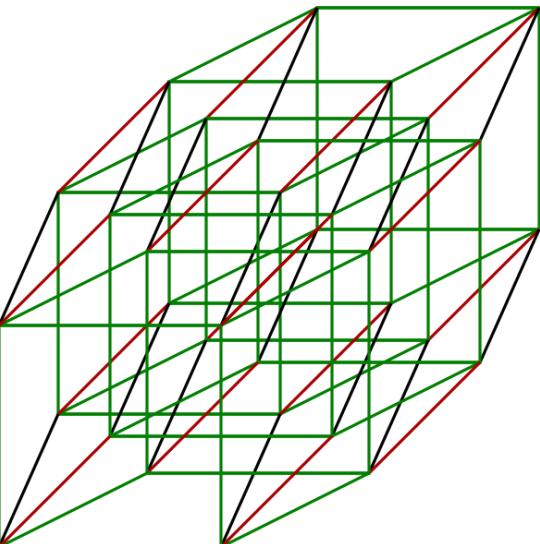
3 Dimensions



4 Dimensions



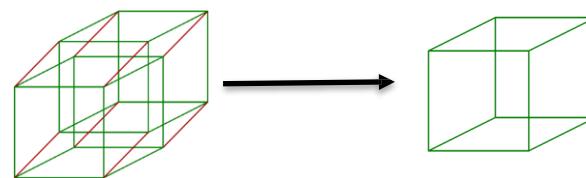
5 Dimensions



- Dimension of an object can be defined as minimum number of coordinates needed to specify any point within it.

Dimensionality Reduction

- Process of deriving a set of degrees of freedom which can be used to reproduce most of the variability of a data set.
- **Goal:** To reduce dimensions by removing redundant and dependent features.
- **How?** By transforming features from higher dimensional space to a lower dimensional space.

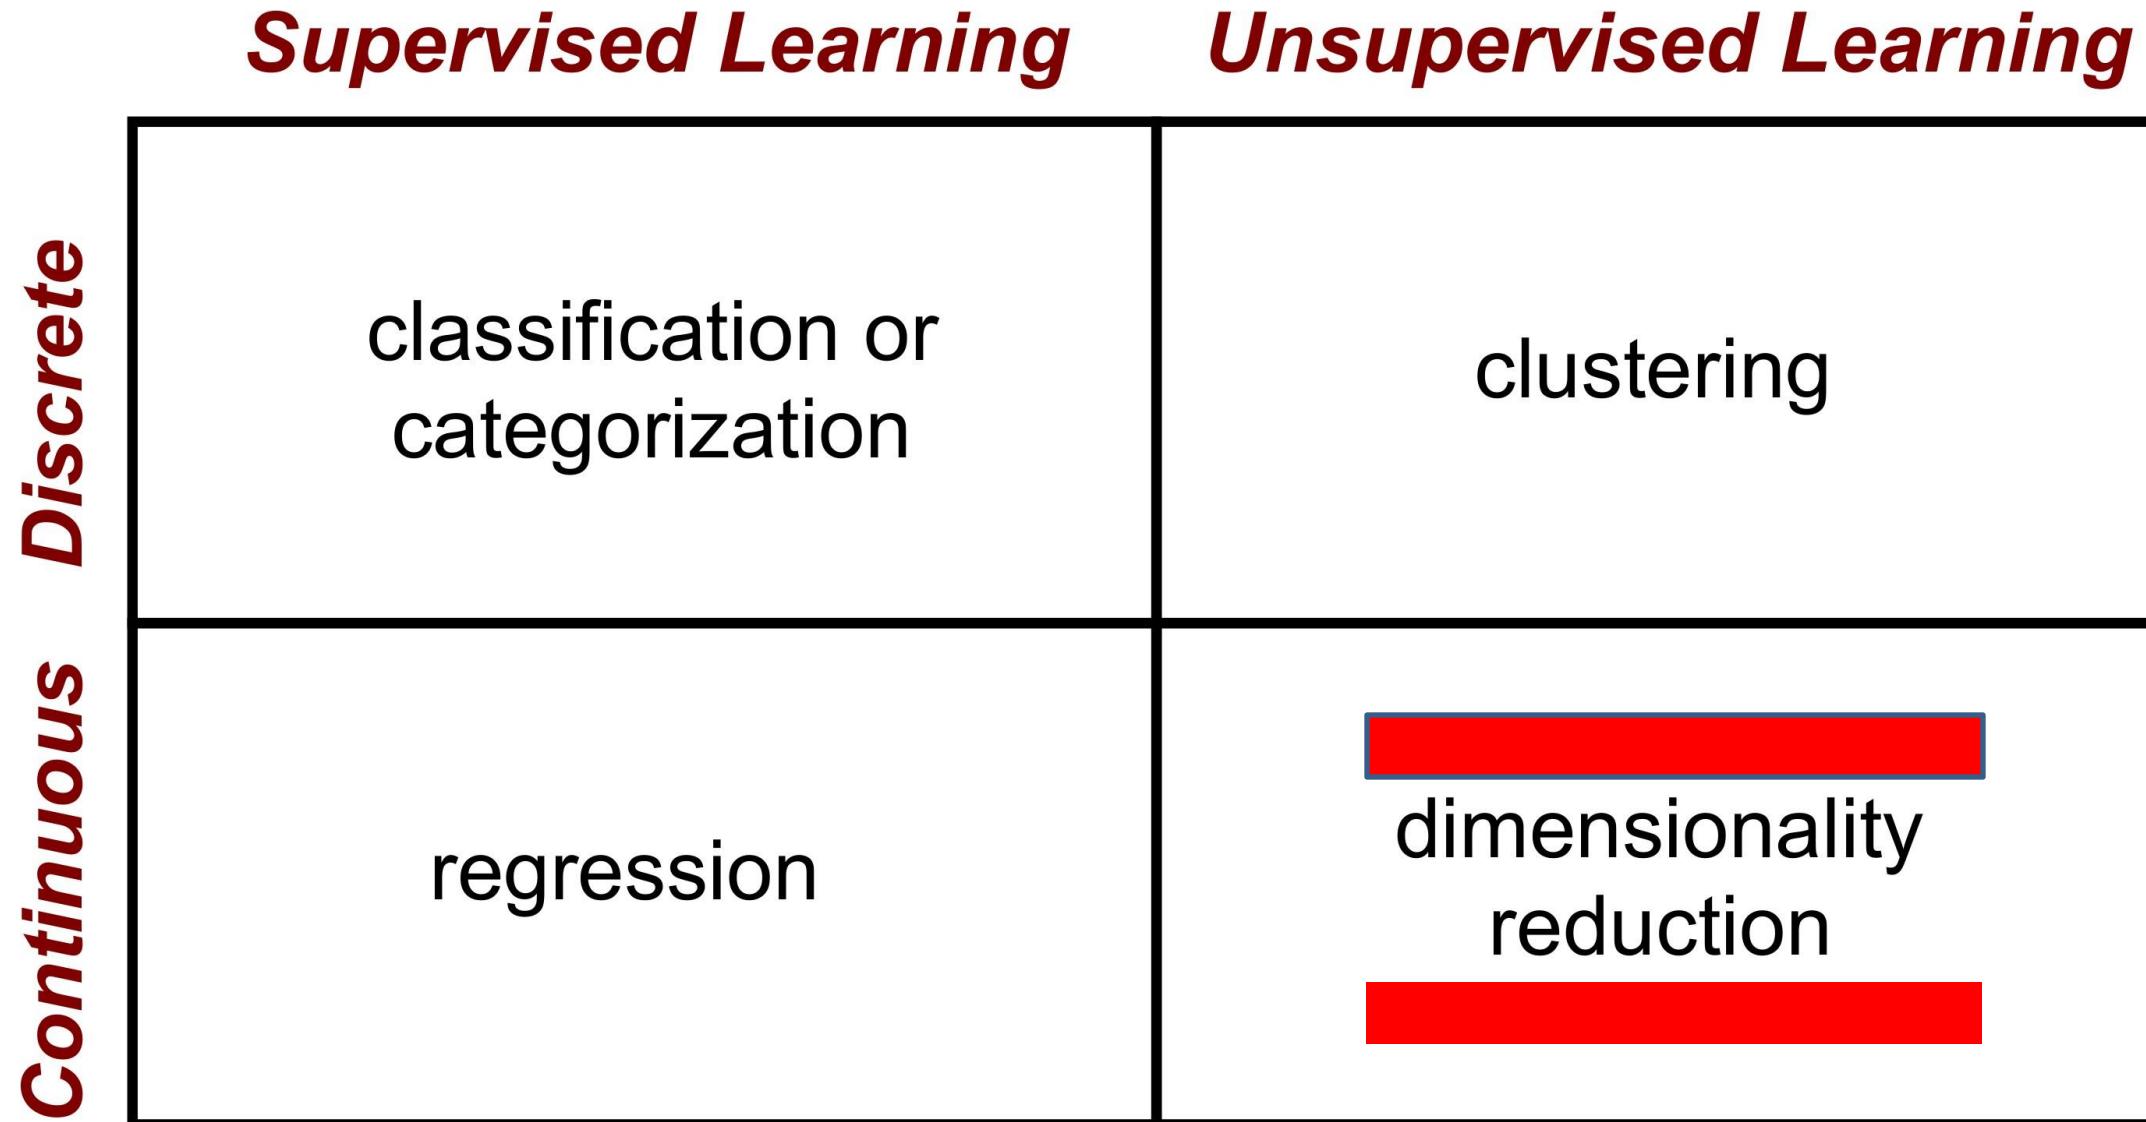


- **Approaches:** Unsupervised and Supervised

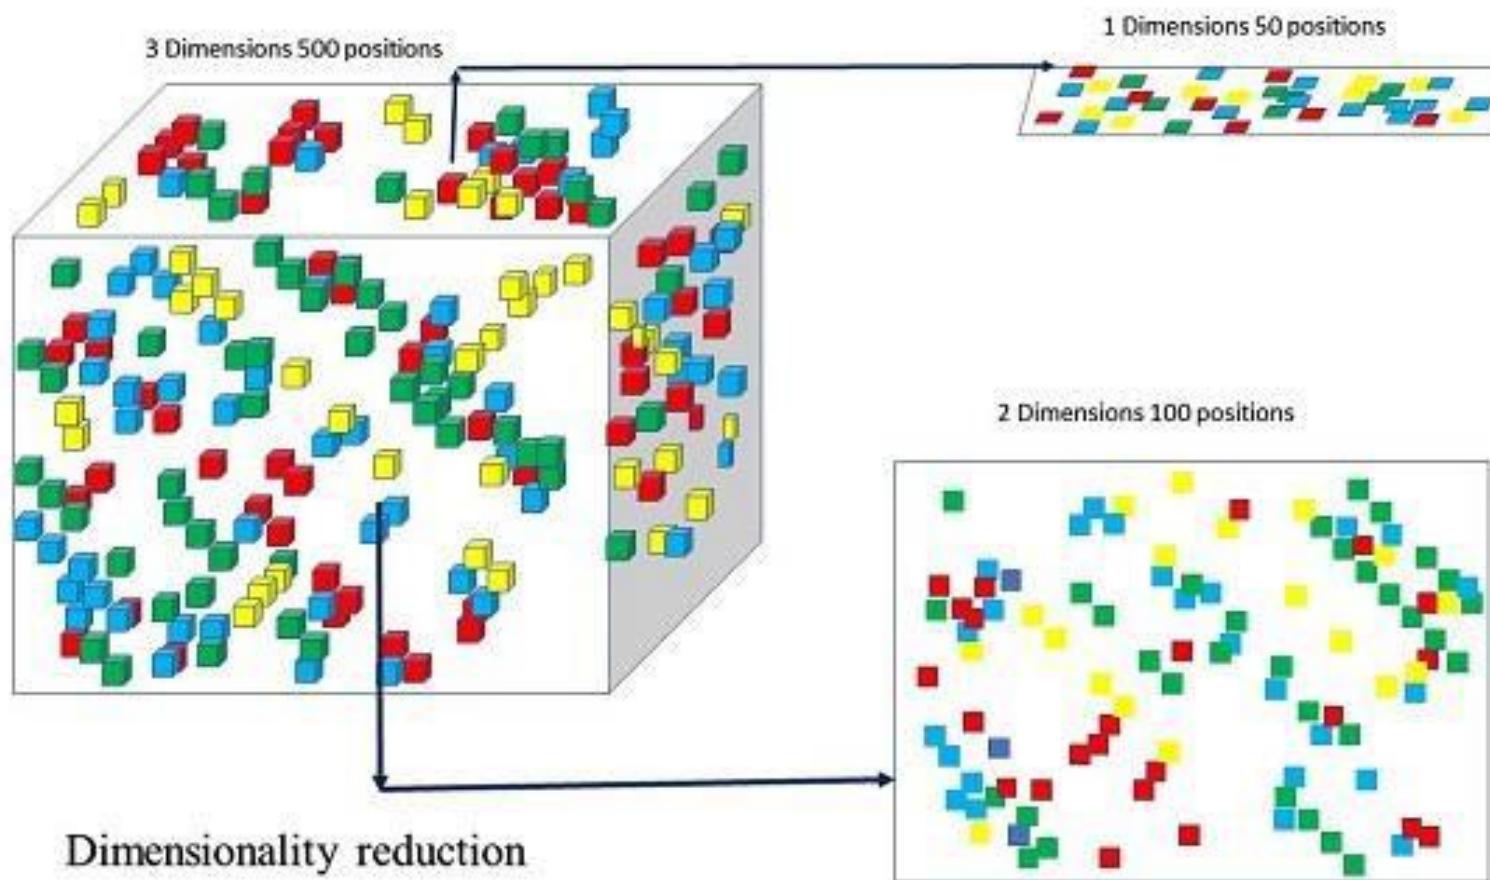
Dimensionality Reduction

- Approaches: Unsupervised and Supervised
 - Unsupervised where no need for labeling classes of data.
 - Independent Component Analysis (ICA)
 - Non-negative Matrix Factorization (NMF)
 - **Principal Component Analysis (PCA)**
 - Ideal for visualization and noise removal.
 - Supervised where class labels are considered.
 - Mixture Discriminant Analysis (MDA)
 - Linear Discriminant Analysis (LDA)
 - Ideal for biometrics, Bioinformatics, and chemistry.

In a nutshell



Dimensionality Reduction



- Application:
 - Machine learning
 - Data mining
 - Bioinformatics
 - Biometric
 - Information Retrieval

Example of multidimensional cube. Perform slicing operations to separate dimensions.

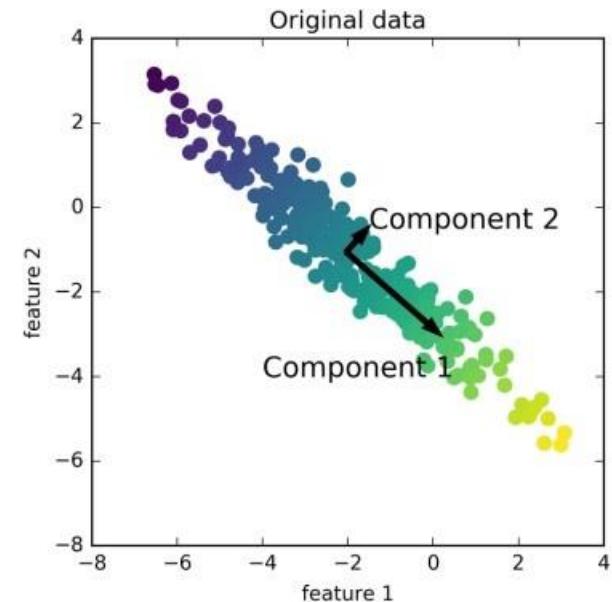
Principal Component Analysis (PCA)

Principal Component Analysis (PCA)

- PCA is a
 - popular technique for dimensionality reduction.
 - "classical" approach only characterize **linear sub-spaces** in data.
- Involves a dataset with observations on numerical variables.
 - An **exploratory** data analysis tool.
 - A **simple, non-parametric** method of extracting relevant information from data sets.

Principal Component Analysis (PCA)

- PCA reduce dimension by exposing underlying information in data sets.
 - An **unsupervised** approach.
 - Explain most of **variability** in data with a **smaller number** of variables.
 - Identifies **axis** that accounts for largest amount of **variance** in training set.
- PCA rotates dataset to get statistically uncorrelated features.
- Rotation followed by selection of subset of new features, based on their importance.



Principal Component Analysis (PCA)

- PCA is influenced by magnitude of each variable.
- First PC (PC_1) of data set is **linear** combination of features that has the largest **variance**.
- Second PC (PC_2) is **linear** combination of data sets that has **maximal variance** out of all linear combinations that are uncorrelated with PC_1 .
- PCs are uncorrelated with each other.
- PCs form a basis of new space (true for irrespective of dimensions).

```
from sklearn.decomposition import PCA  
  
pca = PCA(n_components = 2)
```

Principal Component Analysis (PCA)

- Primary PCA calculation steps:
 - Calculate **covariance** matrix.
 - Calculate **ordered eigenvalues** and **eigenvalues** of the matrix.
- Application:
 - Neuroscience
 - Computer graphics
 - Atmospheric science.

Principal Components (PCs)

- Overall PC calculation process:
 - For each PC:
 - PCA finds a zero-centered **unit vector** pointing in the direction of PC.
 - Direction of unit vectors returned by PCA is not stable*.
 - If you perturb training set slightly and run PCA again
 - Unit vectors may point in opposite direction as original vectors.
 - Still, they will lie on same axes.

Principal Component Analysis (PCA)

- How did we know to use two PCs?
 - Proportion of variance explained (**PVE**) measures exactly what percentage of variance should retain in these PCs.
 - PVE of m^{th} PC is calculated roughly by:
 - $(\text{taking the } m^{\text{th}} \text{ eigenvalue}) / (\text{number of PCs})$.
- How many PCs to Use?
 - A general $n(\text{observations}) \times p(\text{variables})$ data matrix X , there are up to **min($n-1, p$)** PCs.
- No fixed method to determine how many PCs to use.

Demo_PCA_2.ipynb

```
# Access the first principal component (PC1)
pca = pca.components_[0]

# Get feature names and their corresponding contributions to PC1
feature_contributions = pd.DataFrame(pca, index=X.columns, columns=["PC1"])
print("Features contributing to the first principal component (PC1):")
print(feature_contributions)

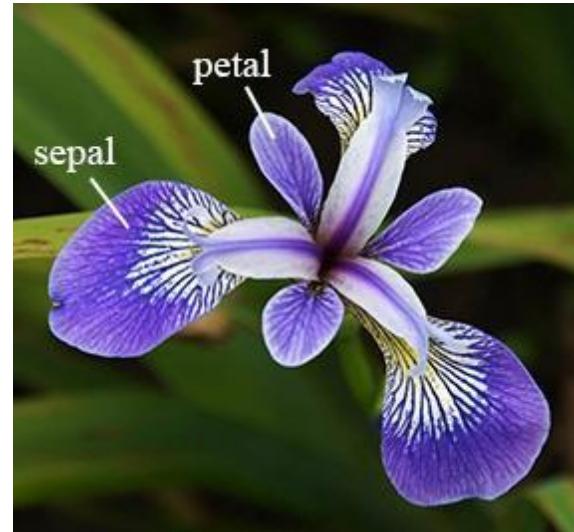
# Sort the features by the absolute value of their contributions to PC1
sorted_features = feature_contributions.abs().sort_values(by="PC1", ascending=False)

# Get the names of the features that contribute to the first principal component
top_features = sorted_features.index.tolist()

# Print out the feature names sorted by contribution
print("Features contributing to the first principal component (PC1):")
print(top_features)

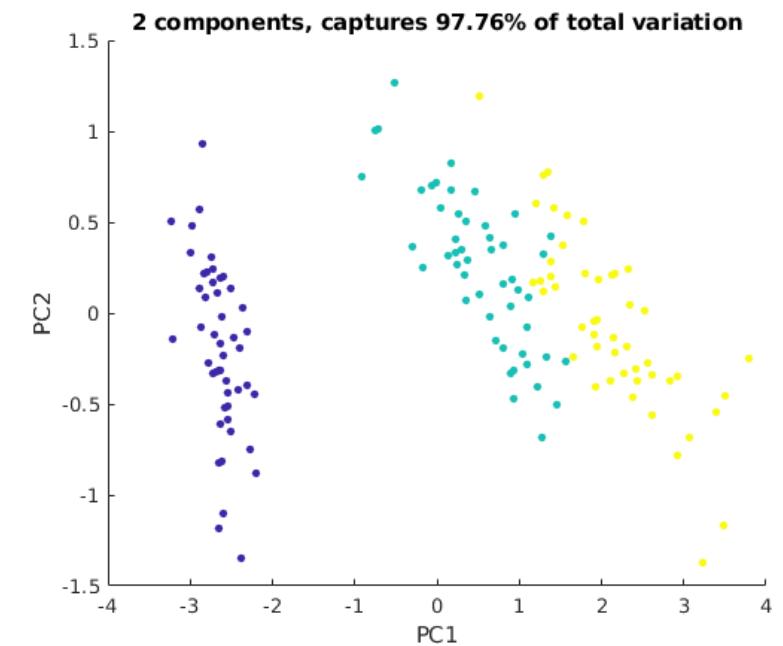
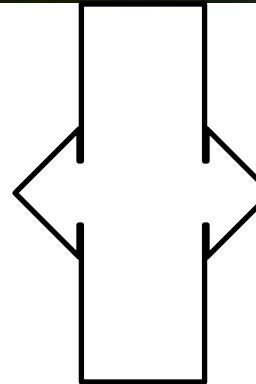
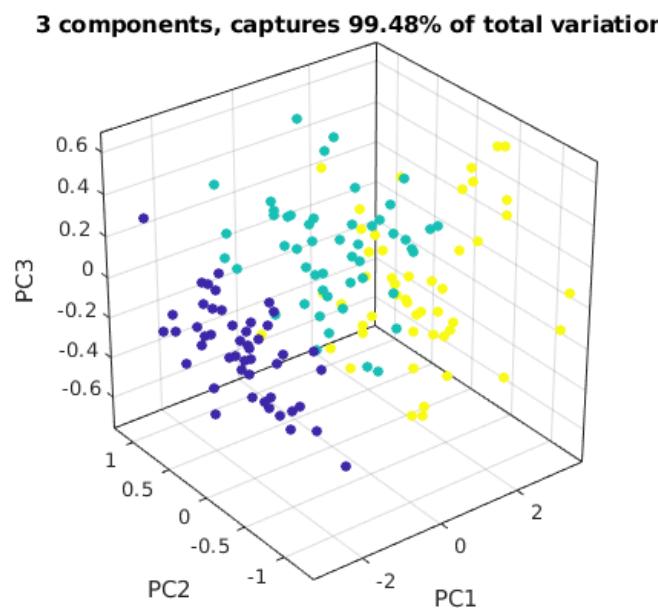
Features contributing to the first principal component (PC1):
PC1
Month      0.002627
Year       0.000108
Patient_Visits  0.999997
Features contributing to the first principal component (PC1):
['Patient_Visits', 'Month', 'Year']
```

PCA Example: Iris data



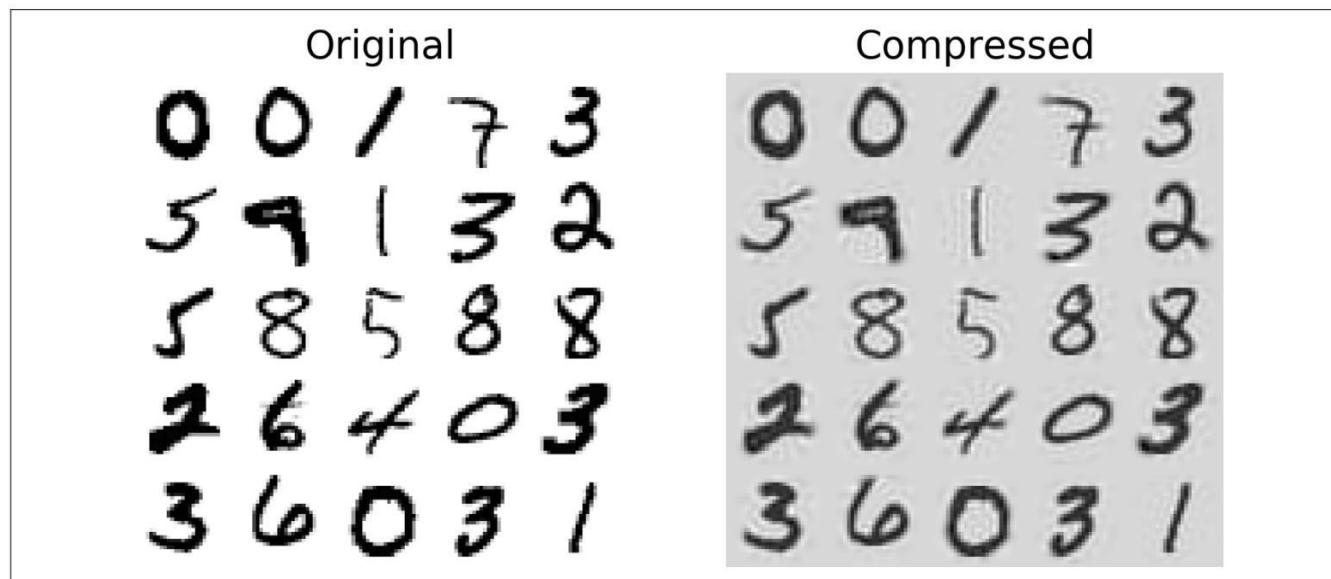
- Total 150 iris flowers.
- Measurement: sepal length, sepal width, petal length, petal width

Flowers belong to three different species: blue, green, yellow



PCA Example: MNIST Compression

- Apply PCA to MNIST dataset while preserving **95%** of its variance.
- Each instance have just over **150 features**, instead of the original **784 features**.
- Dataset is now less than **20%** of its original size.



```
pca = PCA(n_components = 154)
X_reduced = pca.fit_transform(X_train)
X_recovered = pca.inverse_transform(X_reduced)
```

MNIST compression that preserves 95% of the variance

PCA Types

- Types of PCA:
 1. **Randomized PCA** quickly finds an approximation of the first d principal components.
 - Issue: whole training set need to fit in memory.
 2. **Incremental PCA** (IPCA) splits the training set into mini-batches and feed an IPCA algorithm one mini-batch at a time.
 3. **Kernel PCA** helps to perform complex nonlinear projections for dimensionality reduction.

Incremental & Kernel PCA

- IPCA finds a similar projection of data to PCA by processing few samples at a time.
 - Stores estimates of component and noise variances to update variance ratio incrementally.
-
- Kernel PCA can find a projection of the data that makes data linearly separable.
 - Kernel PCA supports both transform and inverse transform.

Singular Value Decomposition (SVD)

[Not in Lecture Plan]

Singular Value Decomposition (SVD)

- SVD is a method for
 - transforming correlated variables into a set of uncorrelated set
 - Exposes various relationships among original data items.
 - Identifying and ordering dimensions along which data points exhibit most variation.
 - Perform data reduction.
 - Power Method is used to compute SVD.

Singular Value Decomposition (SVD)

- Basic steps of SVD:
 - Consider a high dimensional, highly variable set of data points.
 - Reduce it to a lower dimensional space that exposes **substructure** of original data (more clearly).
 - Orders it from most variation to least.

Singular Value Decomposition (SVD)

- SVD decomposes any matrix (A) into three matrixes with special properties and whose multiplication gives back the same matrix: $A_{mn} = U_{mm}S_{mn}V^T_{nn}$
 - U : Orthogonal matrix
 - S : Diagonal matrix
 - V^T : Transpose of an orthogonal matrix
- Note: $U^TU = I$ (Identity Matrix), $V^TV = I$
 - Columns of U are orthonormal eigenvectors of AA^T
 - Columns of V are orthonormal eigenvectors of A^TA
 - S is a diagonal matrix containing square roots of eigenvalues from U or V in descending order.

Singular Value Decomposition (SVD)

- SVD can also be represented as: $A = \sum_{i=1}^r \sigma_i u_i v_i^T$
- Let A be a $n \times d$ matrix with singular vectors v_1, v_2, \dots, v_r and corresponding singular values $\sigma_1, \sigma_2, \dots, \sigma_r$.
- Then $u_i = (1/\sigma_i)^* A v_i$, for $i = 1, 2, \dots, r$, are left singular vectors. Where $\text{Rank}(A) = r$.
- A can be decomposed into a sum of rank one matrices as:
- $A = U \Sigma V^T$

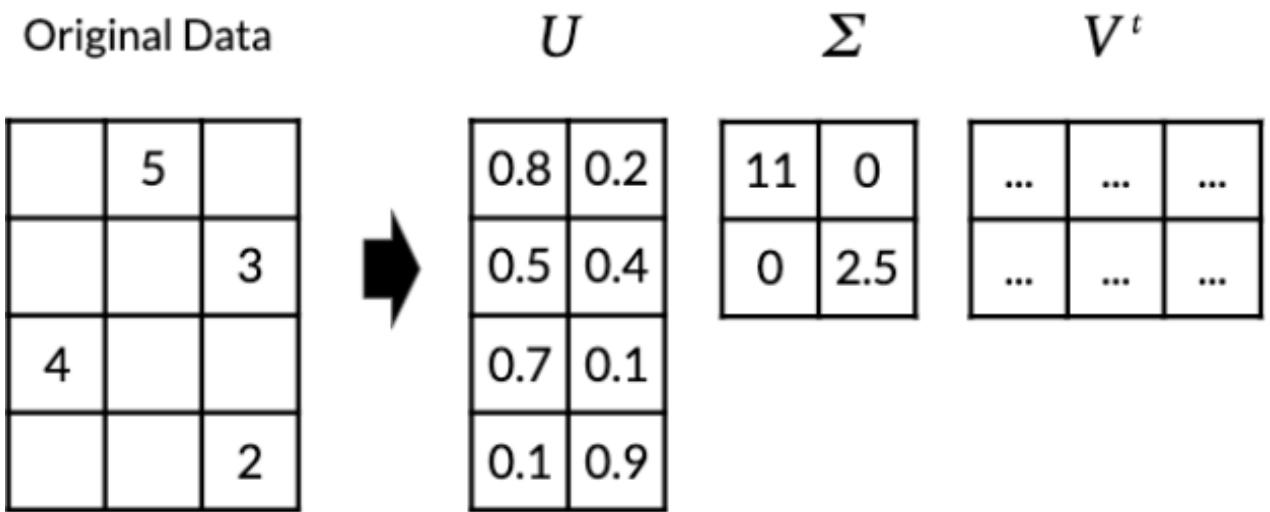
OR

$$A = \sum_{i=1}^r \sigma_i u_i v_i^T$$

$U \in \mathbb{R}^{m \times r}$ has orthonormal columns; $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_r)$, where $\sigma_1 \geq \dots \geq \sigma_r > 0$; $V \in \mathbb{R}^{n \times r}$ has orthonormal columns.

Singular Value Decomposition (SVD)

```
from sklearn.decomposition import TruncatedSVD  
import numpy as np  
  
np.random.seed(0)  
X = np.random.rand(100, 100)  
  
# four components  
svd = TruncatedSVD(n_components=4, n_iter=10, random_state=5)  
  
U = svd.fit_transform(X)  
Sigma = np.diag(svd.singular_values_)  
V = svd.components_  
  
# in case we want to do the multiplication  
  
# U x Sigma  
U_x_Sigma = np.dot(U, Sigma)  
  
# (U x Sigma) x V  
U_Sigma_V = np.dot(U_x_Sigma , V)
```



Application of SVD: Compression

- For many images, k much smaller than n can be used to reconstruct the image provided that a very low-resolution version of the image is sufficient.
- Thus, one could use SVD as a compression method.

Singular Values: 478



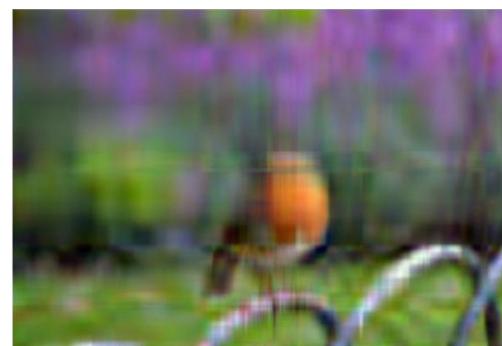
Singular Values: 60



Singular Values: 25



Singular Values: 10



Decision Tree

Decision Trees

- Decision Trees can perform both classification and regression* tasks, and multi-output tasks.
- Decision Trees are used
 - for inductive inference.
 - for approximating discrete-valued target functions.
- Decision Trees are
 - robust to noisy data and can learn disjunctive expressions.
 - capable of fitting complex datasets.
 - capable of leading to a decision via if/else questions.

* estimate the relationships between a dependent variable ('outcome variable') and one or more independent variables.

Decision Trees

- Decision tree searches a completely expressive hypothesis.
 1. Represent a disjunction of conjunctions of constraints on attribute values of instances.
 2. Each path from tree root to a leaf corresponds to a conjunction of attribute tests
 3. Tree itself is a disjunction of these conjunctions.
- Solves classify problems.
 - Learned function is represented by a decision tree.

When to Consider Decision Trees?

- Instances are represented by attribute-value pairs.
- Target function has discrete output values (two or more) (e.g., Boolean classification).
- Disjunctive descriptions may be required (represent disjunctive expressions).
- Training data may contain errors (robust to errors) or missing values (can handle).

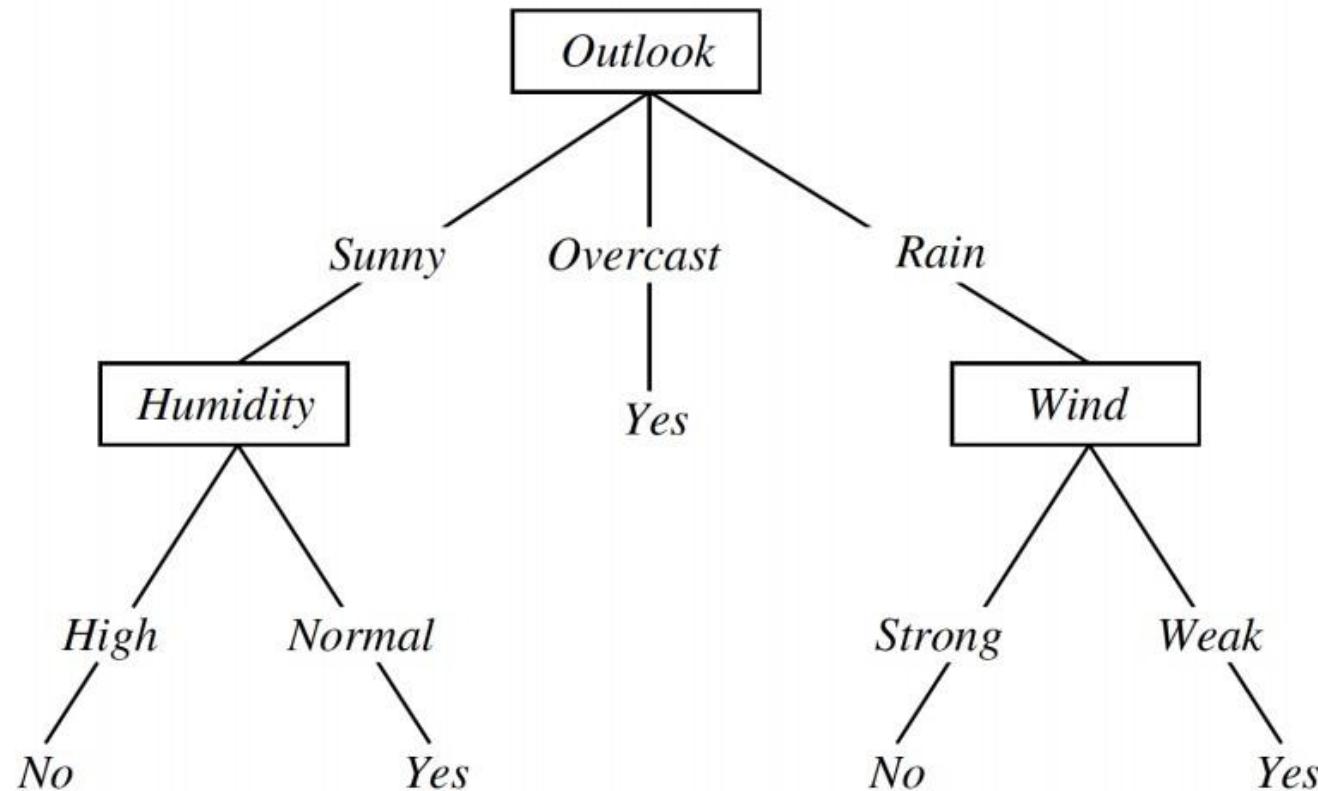
Decision Trees

- Decision trees **classify instances by sorting from** root to leaf node.
- Each node of tree specifies a test of some **attribute** of instance.
- Each branch descending from a node corresponds to one of possible values for attribute.
- Each leaf node assigns a classification.

Example

$(\text{Outlook} = \text{Sunny} \wedge \text{Humidity} = \text{Normal}) \vee (\text{Outlook} = \text{Overcast}) \vee (\text{Outlook} = \text{Rain} \wedge \text{Wind} = \text{Weak})$

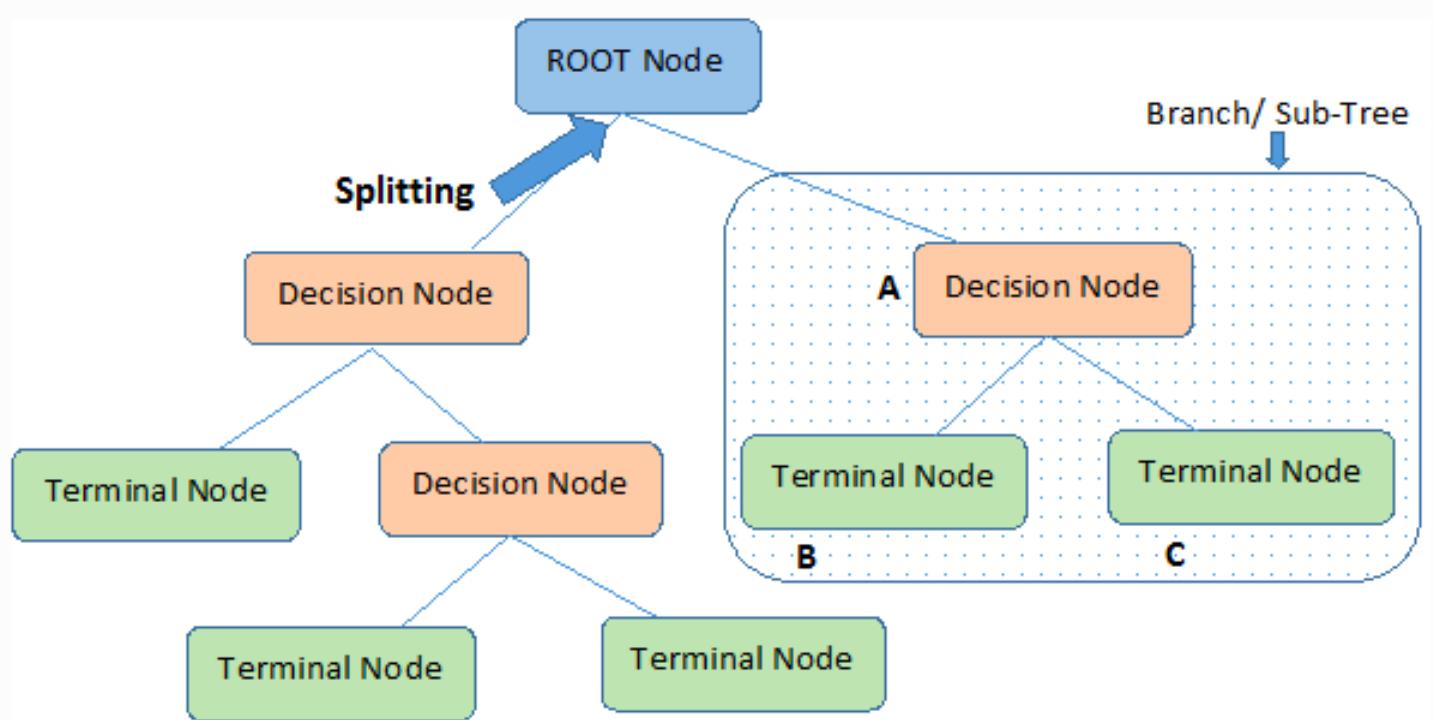
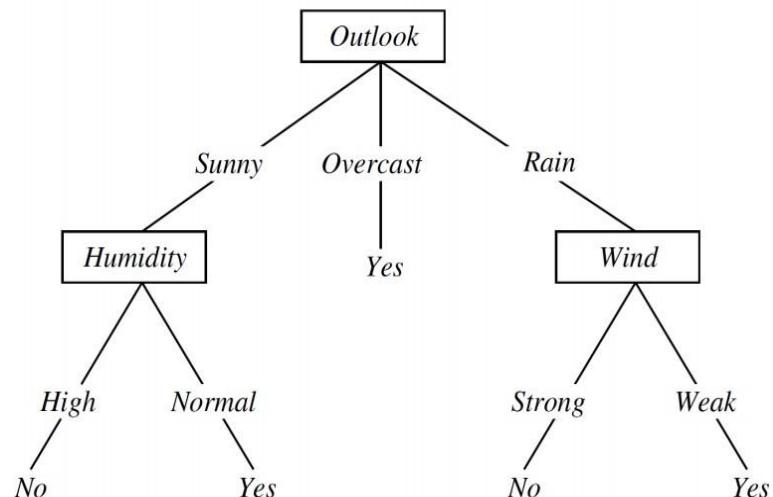
$(\text{Outlook} = \text{Sunny}, \text{Humidity} = \text{High}, \text{Wind} = \text{Strong}) \Rightarrow \text{negative instance.}$



Example

$(\text{Outlook} = \text{Sunny} \wedge \text{Humidity} = \text{Normal}) \vee (\text{Outlook} = \text{Overcast}) \vee (\text{Outlook} = \text{Rain} \wedge \text{Wind} = \text{Weak})$

(Outlook=Sunny, Humidity=High, Wind=Strong) ==> negative instance.

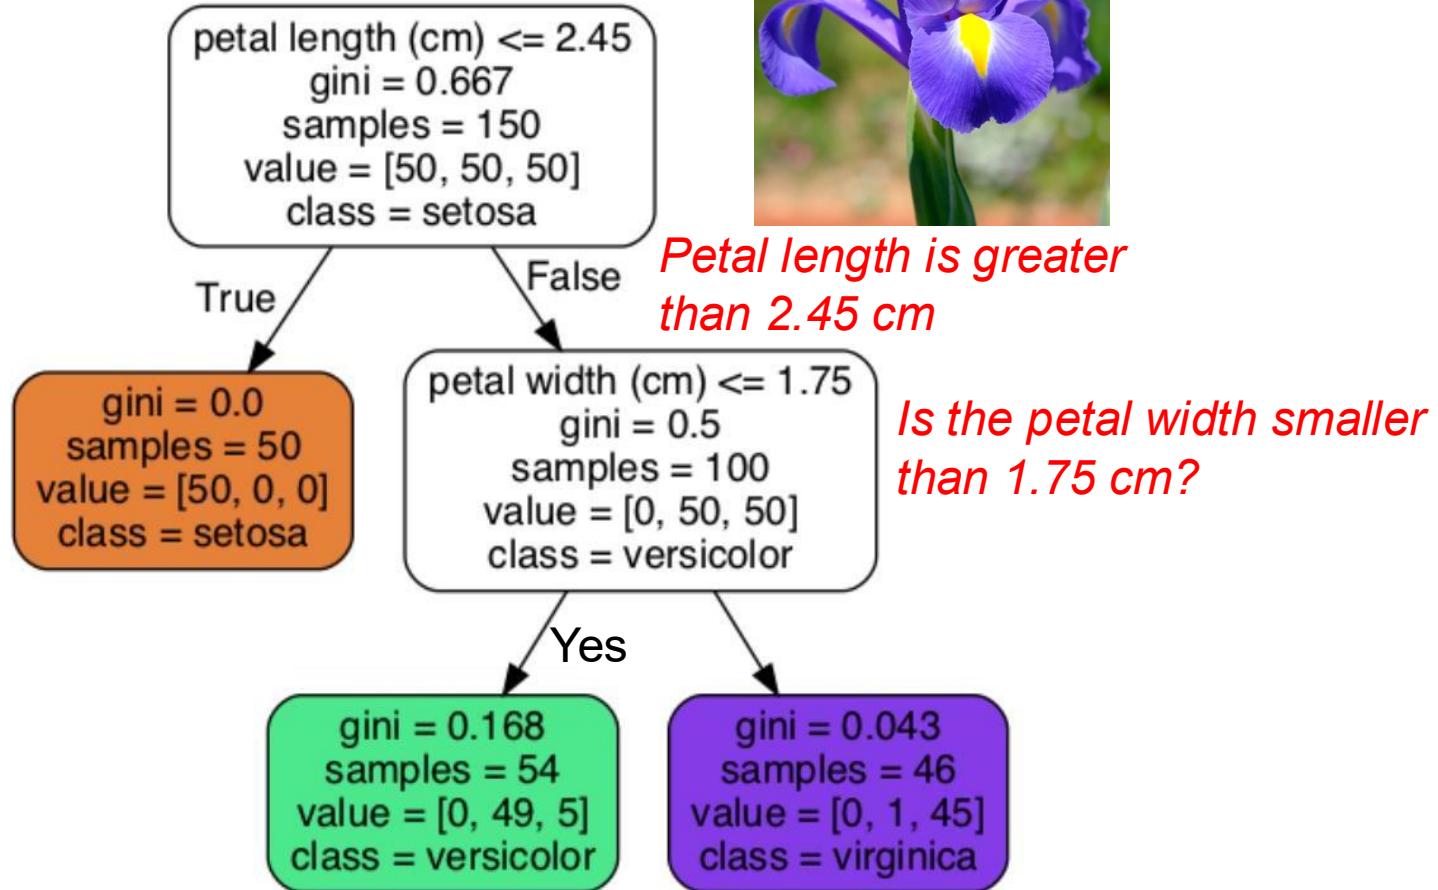


Note:- A is parent node of B and C.

Training and Visualizing a Decision Tree

```
from sklearn.datasets import load_iris  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.tree import export_graphviz  
  
iris = load_iris()  
  
X = iris.data[:, 2:] # petal length and width  
y = iris.target  
  
  
tree_clf = DecisionTreeClassifier(max_depth=2)  
tree_clf.fit(X, y)  
  
export_graphviz(  
    tree_clf,  
    out_file=image_path("iris_tree.dot"),  
    feature_names=iris.feature_names[2:],  
    class_names=iris.target_names,  
    rounded=True,  
    filled=True  
)
```

Q. You want to find classify an iris flower



Petal length is greater than 2.45 cm

Is the petal width smaller than 1.75 cm?

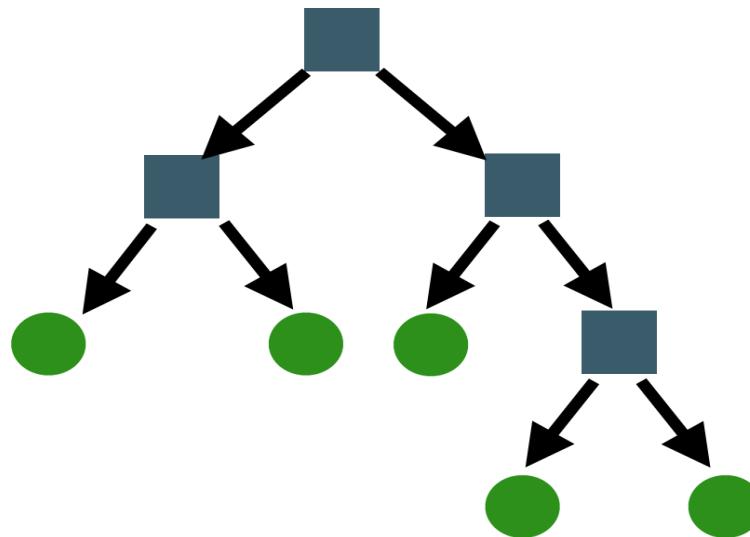
*Require very little data preparation.
Don't require feature scaling.*

Learning Decision Trees

- Given a set of training data in form of examples (e.g., rating), construct questions that you can ask.
- Learning is about searching for the “**best**” tree to describe data.
- We could enumerate all possible trees, and evaluate each
 - How many trees are there given 3 features? too many!
 - It is computational infeasible to consider all trees, so **decision trees must be built greedily by asking important question.**
- Each node represents a question that split data
 - Learning a decision tree amounts to **choosing what internal nodes should be.**

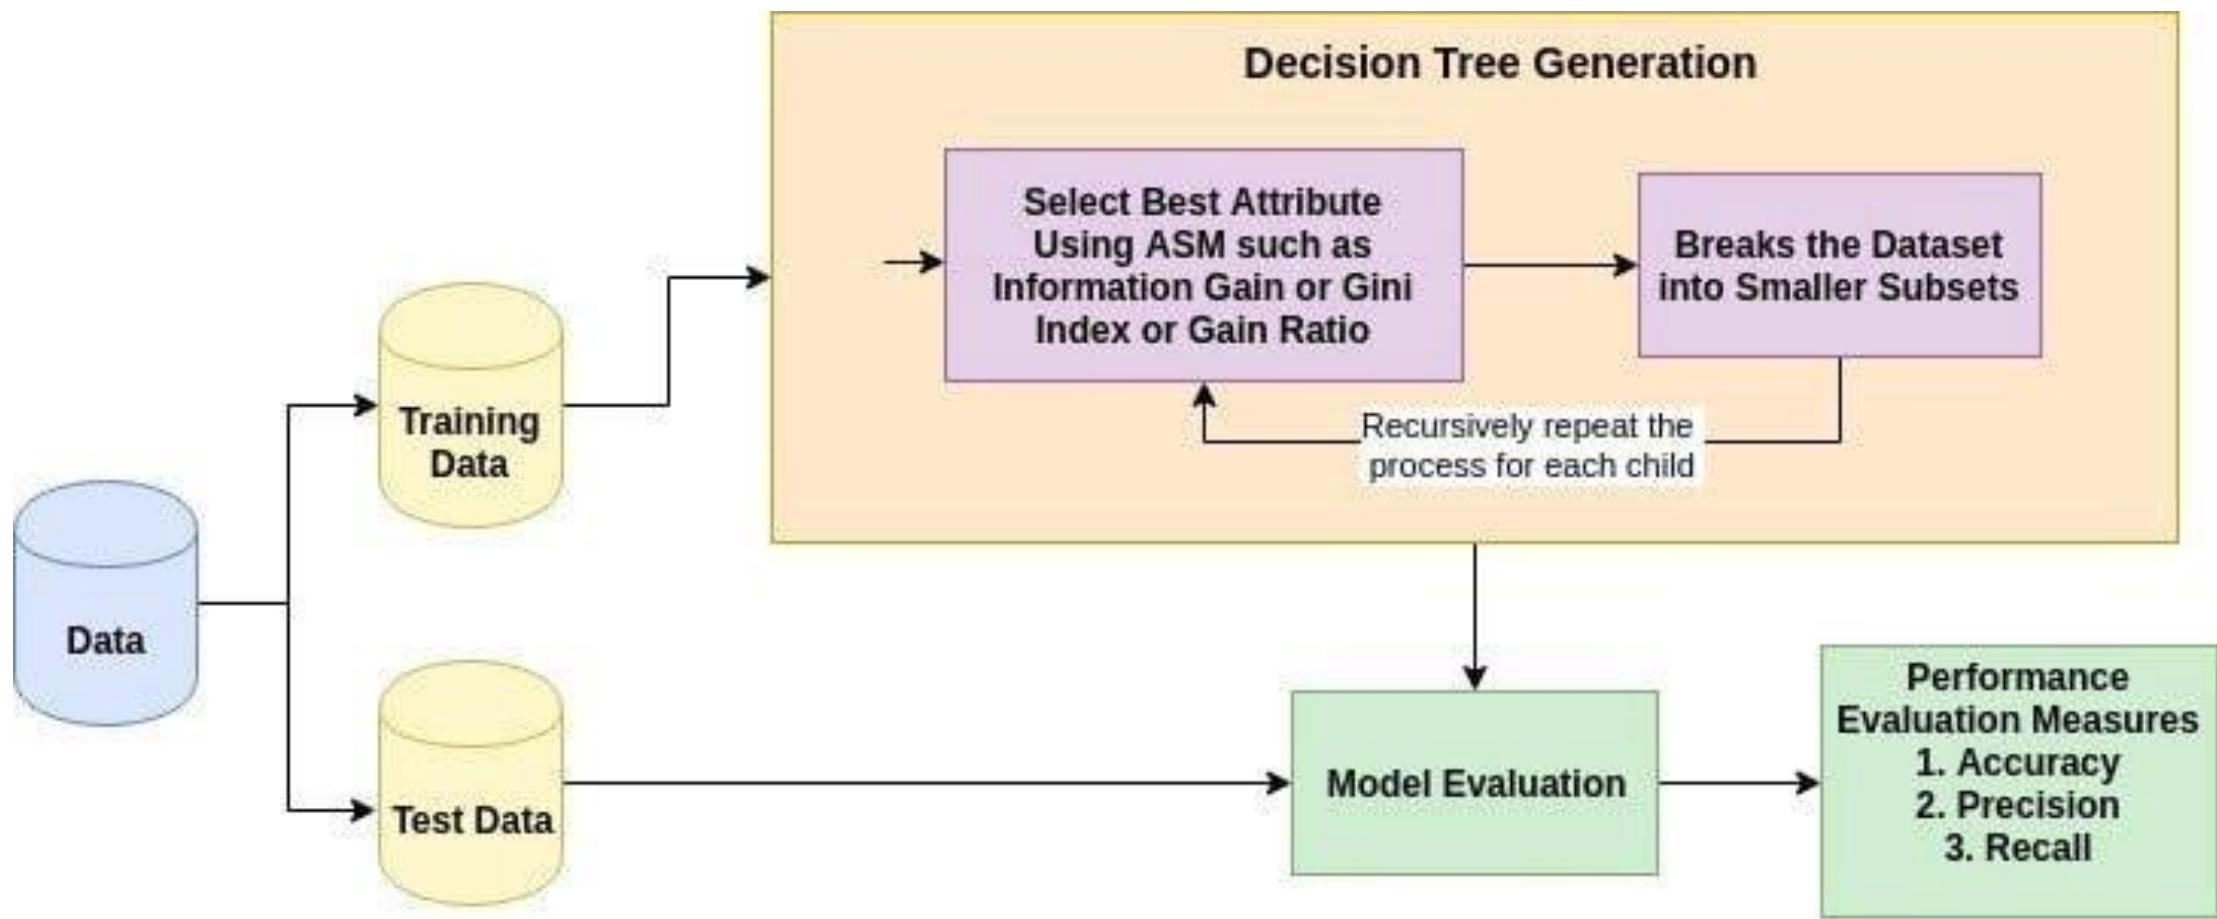
Decision Tree Training Algorithm

- Given a set of labeled training instances:
 - If all the training instances have same class, create a **leaf** with that class label and exit.
 - Else Pick the best, test to split the data on.
 - Split the training set according to the value of the outcome of the test.
 - Recursively repeat step 1-3 on each subset of the training data.



Decision Trees

ASM: Attribute Selection Measure



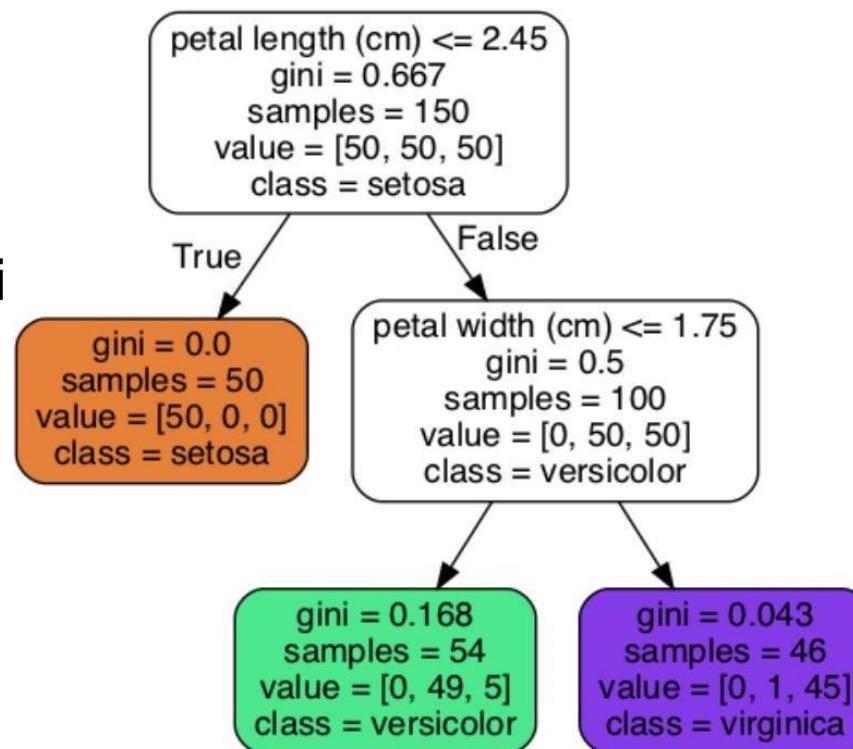
Impurity: Gini and entropy

- Gini Index is to divide decision tree.
- Gini Index (Impurity) calculates the likelihood that somehow a randomly picked instance would be erroneously cataloged.
- Works on **categorical variables** & does **not work with continuous targets**.
- Results in “success” /1 or “failure”/0.
- Gini used to find which attribute holds the maximum information about a class.
 - Enhances accuracy and reliability.

$$G_i = 1 - \sum_{k=1}^n p_{i,k}^2$$

$p_{i,k}$: ratio of instances of class k in node i

Range: 0 (perfectly pure) to 0.5 (maximum impurity, for a binary classification).



Gini impurity at the root node can be more than 0.5.

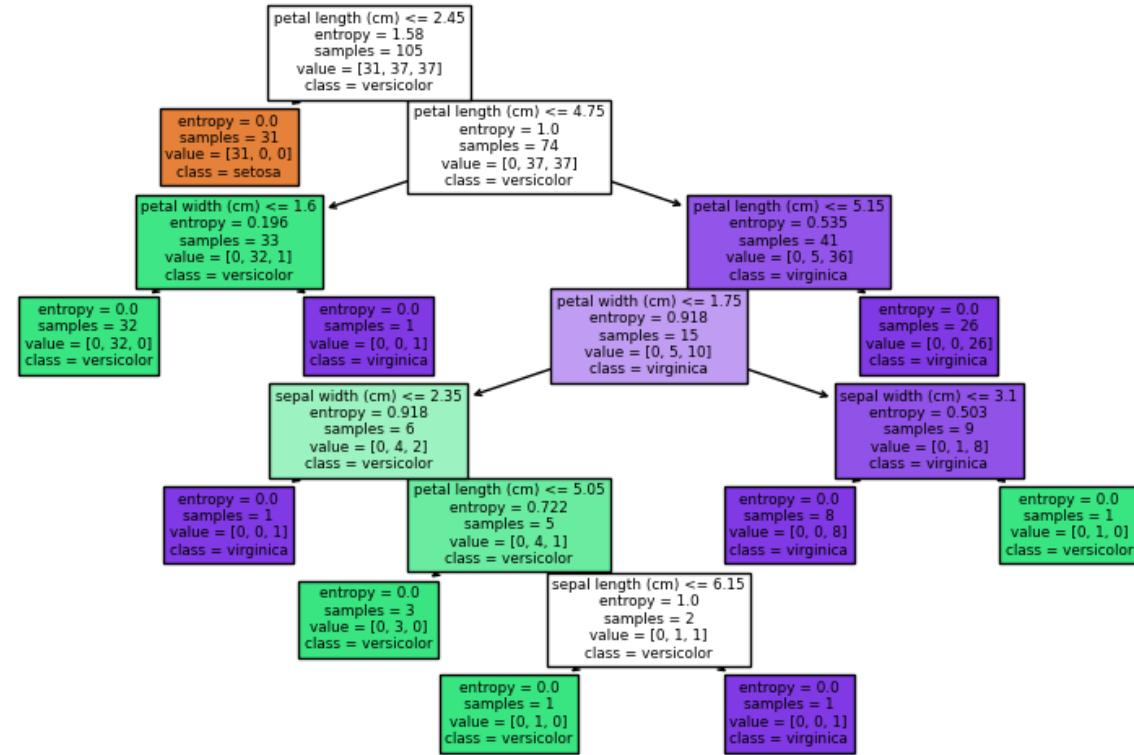
- It means the root node is evenly split among the different classes (see the value = [50, 50, 50]), making it harder to classify the data points correctly without further splitting that's what we did.

Impurity: Gini and Entropy

- Indicates disorder.
- Used as an impurity measure.
 - a set's entropy is zero when it contains instances of only one class (see side pic=>).
- Reduction of entropy is called an **information gain**.
 - The difference between a parent node's entropy and the weighted sum of its child node entropies.

$$H_i = - \sum_{k=1}^n p_{i,k} \log_2 (p_{i,k})$$

$p_{i,k} \neq 0$



Range: 0 (perfectly pure) to $\log_2(n)$ (maximum impurity, n= # of class)

Ensembles: Random Forests

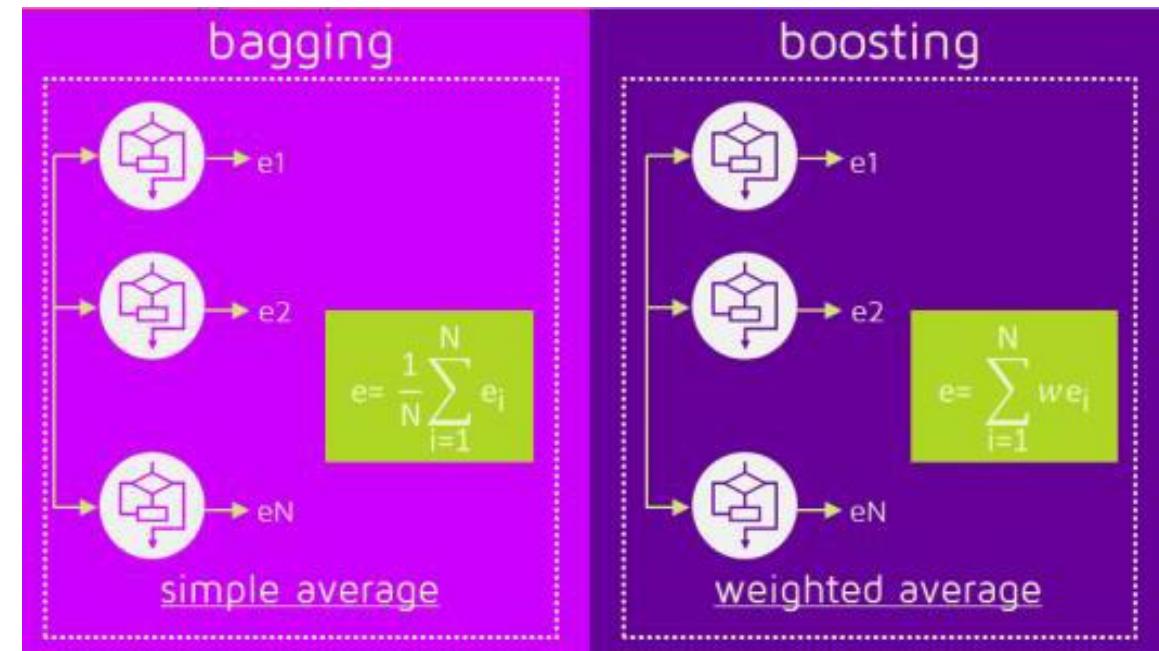
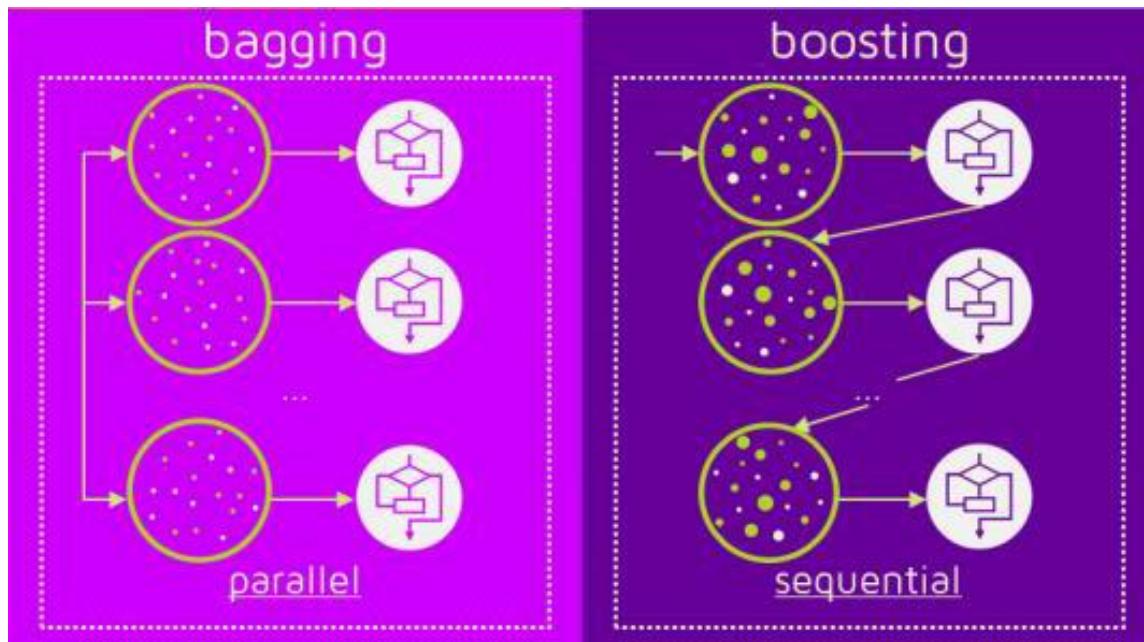
Ensemble Learning

- Noise, variance, and bias are important factors in ML based prediction.
- Ensemble can reduce variance, and bias.
 - is a collection of predictors to give a final prediction.
 - Further classified into **Bagging** and **Boosting**.
- Aggregating several predictors is often better than using just one predictor
 - Several “weak” learners → “strong” learner
- Differences between ensemble learning approaches:
 - **Which** one to combine?
 - **How** to combine?

Bagging and Boosting

- Bagging builds many **independent predictors** and combine them using some model averaging techniques. (e.g. weighted average, majority vote)
 - Random Forest
 - Handle overfitting
- In Boosting, predictors are **not made independently, but sequentially.**
 - Gradient Boosting
 - Can overfit
- Ensemble methods work best when the predictors are as **independent** from each other.

Bagging and Boosting



- Train classifiers using different algorithms.
- Will make different types of errors, improving ensemble's accuracy.

Random Forests

- Random Forest is an ensemble of Decision Trees.
- Trained via bagging method.
- Random Forest Classifier more convenient and optimized for Decision Trees.
- Random Forest classifier with 500 trees:

Lower number of leaf
nodes restricts the depth
of each tree -> reducing
overfitting

```
from sklearn.ensemble import RandomForestClassifier  
  
rnd_clf = RandomForestClassifier(n_estimators=500, max_leaf_nodes=16, n_jobs=-1)  
rnd_clf.fit(X_train, y_train)  
  
y_pred_rf = rnd_clf.predict(X_test)
```

Random Forests

- RandomForestClassifier has all hyperparameters of a DecisionTreeClassifier (mostly).
 - Can control growth of trees.
 - Hyperparameters of a BaggingClassifier to control the ensemble.
- Random Forest algorithm introduces extra randomness when growing trees
 - Searches for best feature among a random subset of features.
 - Results in greater tree diversity → overall better model.

```
bag_clf = BaggingClassifier(  
    DecisionTreeClassifier(splitter="random", max_leaf_nodes=16), n_estimators=500,  
    max_samples=1.0, bootstrap=True, n_jobs=-1)
```

Each tree is trained
using 100% of the available
training data

Each tree is trained on a
randomly sampled subset

Random Forests

- Tree growth in a Random Forest: at each node only a random subset of features is considered for splitting.
- Possible to make trees more random by using random thresholds for each feature.
- Easy to measure relative importance of each feature (weighted average).
- Ex: Trains a RandomForestClassifier on the iris dataset and outputs each feature's importance.

```
from sklearn.datasets import load_iris
iris = load_iris()

rnd_clf = RandomForestClassifier(n_estimators=500, n_jobs=-1)
rnd_clf.fit(iris["data"], iris["target"])
for name, score in zip(iris["feature_names"], rnd_clf.feature_importances_):
    print(name, score)
```

sepal length (cm)	0.112492250999
sepal width (cm)	0.0231192882825
petal length (cm)	0.441030464364
petal width (cm)	0.423357996355

Most important features are petal length (44%) and petal width (42%)