

# Machine Learning and Deep Learning

## Lecture-12

By: Somnath Mazumdar  
Assistant Professor  
[sma.digi@cbs.dk](mailto:sma.digi@cbs.dk)

# Overview

- Autoencoder
- Hyper-parameter Optimization (HPO)

# Autoencoder

# Autoencoder

- Autoencoders are:
  - unsupervised (self-supervised learning)
  - generative models: capable of randomly generating new data that looks very similar to training data.
- ANNs capable of learning dense representations of input data, called latent representations or codings.
- Codings typically have a much lower dimensionality than input data.
- Autoencoders useful **for dimensionality reduction** good for visualization purposes.
- Autoencoders **simply learn to copy their inputs to their outputs.**

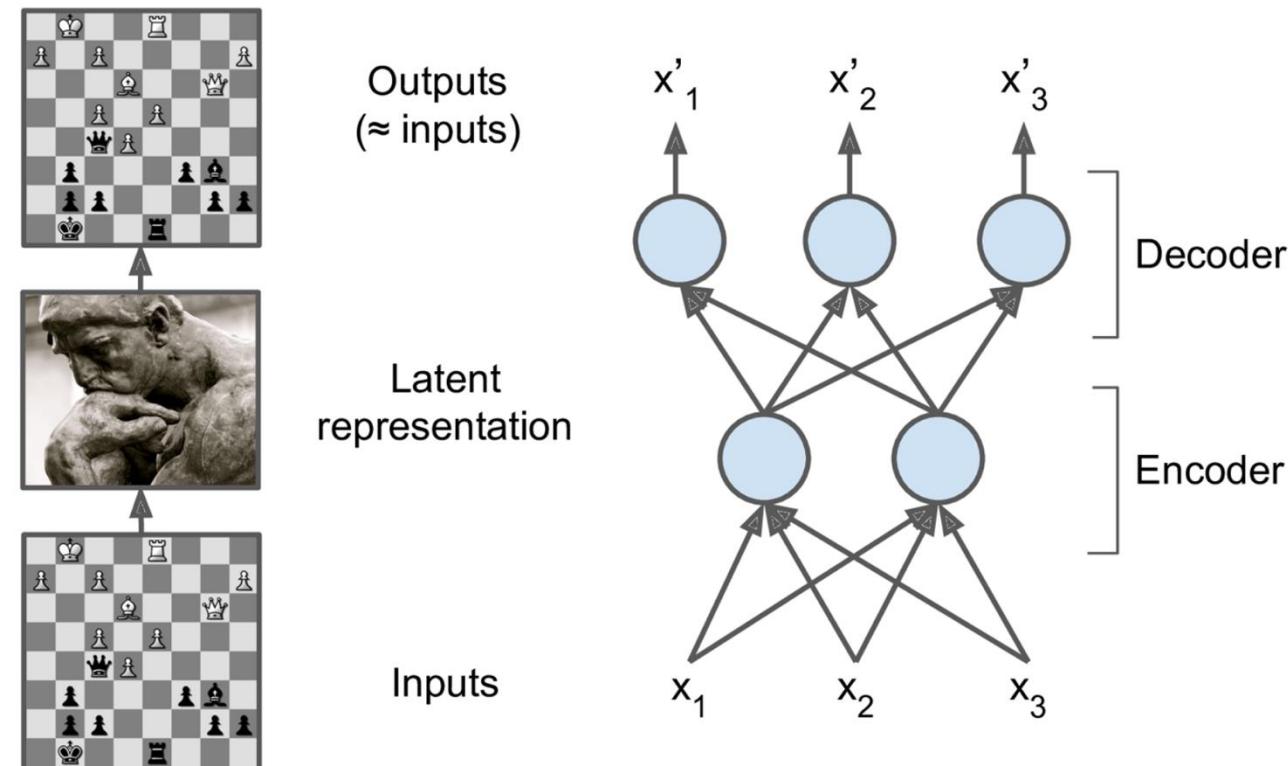
# Autoencoder

- **Codings** are byproducts of autoencoder learning **identity function** under some constraints.
- An autoencoder looks at inputs, converts them to an efficient latent representation, and then output something that (hopefully) looks very close to inputs.
- E.g.,: you can add noise to inputs and train network to recover original inputs.
- Autoencoder composed of two parts:
  - An encoder (or recognition network) that converts inputs to a latent representation.
  - A decoder (or generative network) that converts internal representation to outputs.

# Autoencoder

- Autoencoder has same architecture as MLP except number of neurons in output layer must be equal to number of inputs.

- Here is single hidden layer composed of two neurons (encoder).
- One output layer composed of three neurons (decoder).
- Outputs are often called **reconstructions** because autoencoder tries to reconstruct inputs

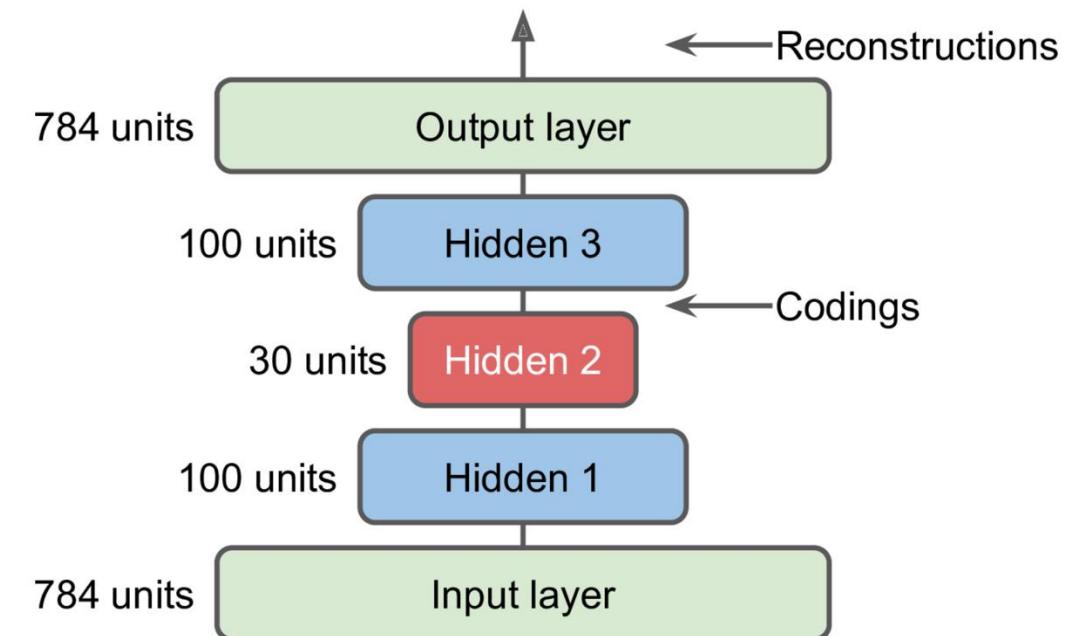


# Autoencoder

- Cost function contains a reconstruction loss that **penalizes** model when reconstructions are different from inputs.
- **Undercomplete autoencoder** where internal representation has a lower dimensionality than input data.
  - An undercomplete autoencoder cannot trivially copy its inputs to codings, and forced to learn most important features in input data (and drop unimportant ones).
- If autoencoder uses only linear activations and cost function is mean squared error (MSE), then it ends up performing Principal Component Analysis (PCA).

# Stacked Autoencoder

- Stacked autoencoders (or deep autoencoders): Autoencoders can have multiple hidden layers.
  - Advantage: Adding more layers helps autoencoder learn more complex codings.
- An autoencoder for MNIST have 784 inputs, followed by a hidden layer with 100 neurons.
- A central hidden (coding) layer of 30 neurons, then another hidden layer with 100 neurons, and an output layer with 784 neurons



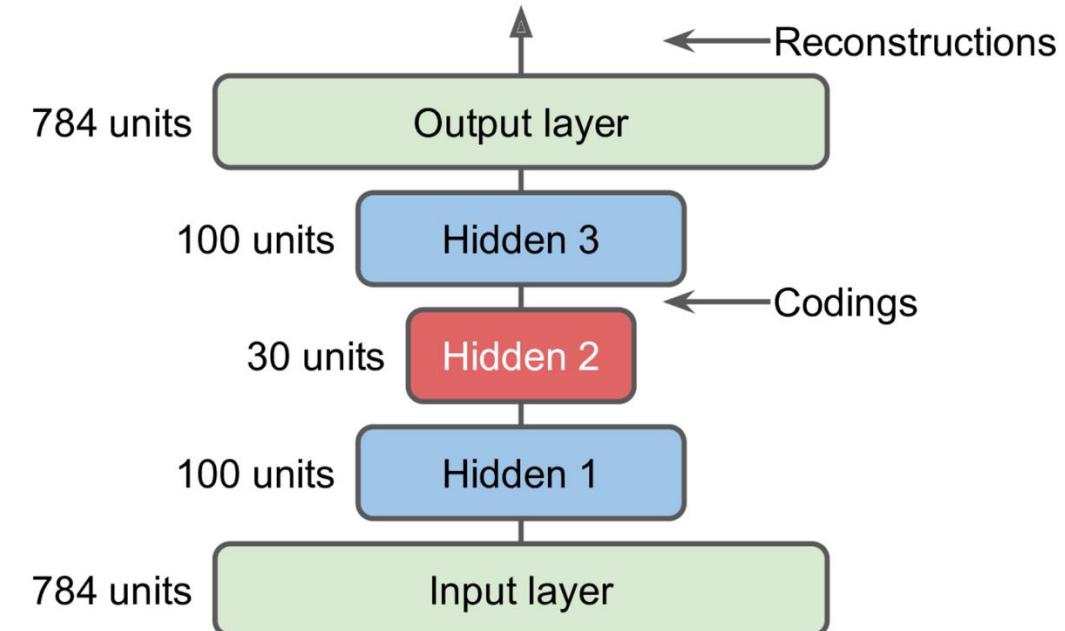
# Stacked Autoencoder

- Stacked autoencoder very much like a regular deep MLP.

```
stacked_encoder = keras.models.Sequential([
keras.layers.Flatten(input_shape=[28, 28]),
keras.layers.Dense(100, activation="selu"),
keras.layers.Dense(30, activation="selu"),
])
```

```
stacked_decoder = keras.models.Sequential([
keras.layers.Dense(100, activation="selu", input_shape=[30]),
keras.layers.Dense(28 * 28, activation="sigmoid"),
keras.layers.Reshape([28, 28])
])
```

```
stacked_ae = keras.models.Sequential([stacked_encoder, stacked_decoder])
stacked_ae.compile(loss="binary_crossentropy",
optimizer=keras.optimizers.SGD(learning_rate=1.5))
history = stacked_ae.fit(X_train, X_train, epochs=10,
validation_data=[X_valid, X_valid])
```



# Convolution Autoencoder

- If you want to build an autoencoder for images, you need to build a **convolutional autoencoder**.
  - Encoder is a regular CNN composed of convolutional layers and pooling layers.
  - It reduces spatial dimensionality of inputs (i.e., height and width) while increasing depth (i.e., number of feature maps).
  - Decoder must do reverse (upscale image and reduce its depth back to original dimensions), and for this use transpose convolutional layers.

# Recurrent & Denoising Autoencoders

- If you want to build an autoencoder for sequences (such as time series or dimensionality reduction), then RNNs may be better suited than dense networks.
- **Recurrent autoencoder:**
  - Encoder is typically a sequence-to-vector RNN which compresses input sequence down to a single vector.
  - Decoder is a vector-to-sequence RNN that does reverse.
- **Denoising Autoencoders:** force autoencoder to learn useful features by adding noise to its inputs, training it to recover original (noise-free).
- Noise can be pure Gaussian noise added to inputs, or it can be randomly switched-off inputs (dropout).
- Autoencoders could also be used for feature extraction\*.

# Recap

- Autoencoders learn to copy their inputs to their outputs by converting inputs into an efficient latent representation and then outputting something close to the original inputs.
  - Structure: An autoencoder is composed of two main parts:
    - Encoder (or recognition network): Converts inputs to a latent representation.
    - Decoder (or generative network): Converts the internal representation back to outputs.
- The cost function of an autoencoder contains a reconstruction loss that penalises the model when the reconstructions differ from the inputs.
- An autoencoder has the same architecture as a Multilayer Perceptron (MLP), except that the number of neurons in the output layer must be equal to the number of inputs.
- A recurrent autoencoder typically has an encoder that is a sequence-to-vector RNN compressing the input sequence into a single vector, and a decoder that is a vector-to-sequence RNN doing the reverse.
- Applications: Autoencoders are useful for dimensionality reduction (good for visualization)

# Hyper-parameter Optimization (HPO)

# Hyper-parameters

- NNs have several hyperparameters: learning rate, weight of regularization.
- “Hyperparameter” regulate design of model.
- Different from more fundamental parameters representing weights of connections in NN.
  - Primary model parameters weights are optimized with backpropagation only.
  - After fixing the hyperparameters either manually or with the use of a tuning phase.
- Hyperparameters should not be tuned using same data (SGD).
  - Portion of data is held out as validation data.
  - Performance of model is tested on validation set with various choices of hyperparameters.
- Ensures not overfit of training data set.

# Challenges Hyper-parameter Optimization

- Function evaluations can be extremely expensive for large models (e.g., in deep learning), complex machine learning pipelines, or large datasets.
- Configuration space is often complex (comprising a mix of continuous, categorical and conditional hyperparameters) and high-dimensional (not always clear which hyperparameters to optimize).
- No access to a gradient of hyperparameter loss function. Other properties of target function often used in classical optimization do not typically apply (such as convexity and smoothness).
- One cannot directly optimize for generalization performance as training datasets are of limited size.

# Hyper-parameters

ML Algorithm	Main HPs	Optional HPs	HPO methods	Libraries
Linear regression	-	-	-	-
Ridge & lasso	alpha	-	BO-GP	Skpot
Logistic regression	penalty, c, solver	-	BO-TPE, SMAC	Hyperopt, SMAC
KNN	n_neighbors	weights, p, algorithm	BOs, Hyperband	Skpot, Hyperopt, SMAC, Hyperband
SVM	C, kernel, epsilon (for SVR)	gamma, coef0, degree	BO-TPE, SMAC, BOHB	Hyperopt, SMAC, BOHB
NB	alpha	-	BO-GP	Skpot
DT	criterion, max_depth, min_samples_split, min_samples_leaf, max_features	splitter, min_weight_fraction_leaf, max_leaf_nodes	GA, PSO, BO-TPE, SMAC, BOHB	TPOT, Optunity, SMAC, BOHB
RF & ET	n_estimators max_depth, criterion, min_samples_split, min_samples_leaf, max_features	splitter, min_weight_fraction_leaf, max_leaf_nodes	GA, PSO, BO-TPE, SMAC, BOHB	TPOT, Optunity, SMAC, BOHB
XGBoost	n_estimators, max_depth, learning_rate, subsample, colsample_bytree, estimators, voting	min_child_weight, gamma, alpha, lambda	GA, PSO, BO-TPE, SMAC, BOHB	TPOT, Optunity, SMAC, BOHB
Voting Bagging	base_estimator, n_estimators	weights max_samples, max_features	GS GS, BOs	Sklearn sklearn, Skpot, Hyperopt, SMAC
AdaBoost	base_estimator, n_estimators, learning_rate	-	BO-TPE, SMAC	Hyperopt, SMAC
Deep learning	number of hidden layers, 'units' per layer, loss, optimizer, Activation, learning_rate, dropout rate, epochs, batch_size, early stop patience	number of frozen layers (if transfer learning is used)	PSO, BOHB	Optunity, BOHB
K-means	n_clusters	init, n_init, max_iter	BOs, Hyperband	Skpot, Hyperopt, SMAC, Hyperband
Hierarchical clustering	n_clusters, distance_threshold	linkage	BOs, Hyperband	Skpot, Hyperopt, SMAC, Hyperband
DBSCAN	eps, min_samples	-	BO-TPE, SMAC, BOHB	Hyperopt, SMAC, BOHB
Gaussian mixture	n_components	covariance_type, max_iter, tol	BO-GP	Skpot
PCA	n_components	svd_solver	BOs, Hyperband	Skpot, Hyperopt, SMAC, Hyperband
LDA	n_components	solver, shrinkage	BOs, Hyperband	Skpot, Hyperopt, SMAC, Hyperband

# Parameter Initialization

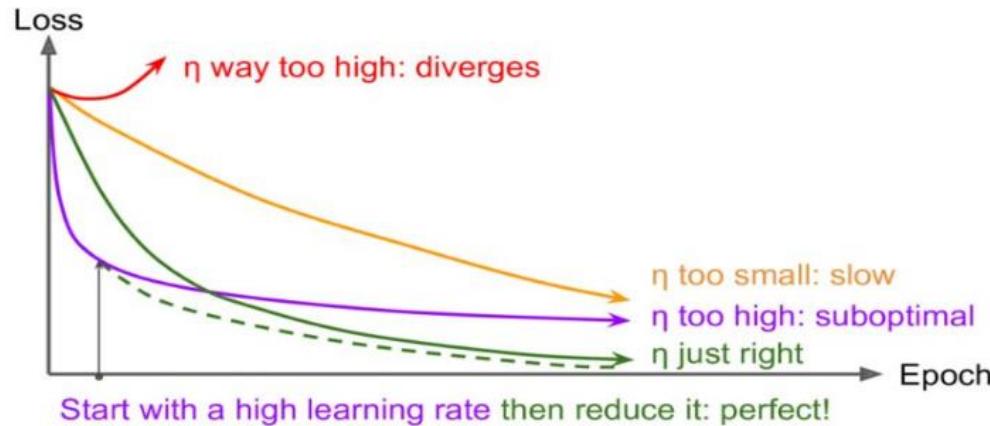
- Parameter initialization is crucial to NN performance:
  - Can't initialize weights in same layer to same value, or they will stay same.
  - Can't initialize weights too large, it will take too long to learn.
- A traditional random initialization:
  - Initialize bias variables to 0.
  - Sample from standard normal, divided by  $10^5(0.00001 * \text{randn})$ .
  - $w = .00001 * \text{randn}(k, 1)$
  - Performing multiple initializations does not seem to be important.
- Popular approach from 10 years ago:
  - Initialize with deep unsupervised model (like “autoencoders”)

# Learning Rate: Setting the Step-Size

- Stochastic gradient is very sensitive to step size in deep models.
- A common approach:
  - Run SG for a while with a fixed step-size.
  - Occasionally measure error and plot progress.
  - If error is not decreasing, decrease step-size.
- Bias step-size multiplier: use bigger step-size for bias variables.
- Momentum: Add term that moves in previous direction where  $\beta^t=0.9$

# Learning Rate Scheduling

- Good learning rate is very important.
  - If you set it much too high, training may diverge (“Gradient Descent”)
  - If you set it too low, training will eventually converge to optimum, at cost of very long time.



- Find a good learning rate by:
  - Training model for a few hundred iterations, exponentially increasing learning rate from a very small value to a very large value.
  - Next look at learning curve and pick a learning rate slightly lower than one at which learning curve starts shooting back up.
  - Reinitialize model and train it with that learning rate.

# Learning Rate Scheduling

- Strategies to reduce learning rate during training called **learning schedules**.
  1. **Power scheduling**: Set learning rate to a function of iteration number  $t$ :  $\eta(t) = \eta_0 / (1+t/s)^c$  . (initial learning rate  $\eta_0$ , power  $c$  (set to 1), steps  $s$  are hyperparameters).
  2. **Exponential scheduling**: Set learning rate to  $\eta(t) = \eta_0 0.1^{t/s}$ . Learning rate will gradually drop by a factor of 10 every  $s$  steps.
  3. **Piecewise constant scheduling**: Use a constant learning rate for a number of epochs (e.g.,  $\eta_0 = 0.1$  for 5 epochs), then smaller learning rate for another number of epochs (e.g.,  $\eta_1 = 0.001$  for 50 epochs), and so on.
  4. **Performance scheduling**: Measure validation error every  $N$  steps (like early stopping) and reduce learning rate by a factor of  $\lambda$  when error stops dropping.

# Hyper-parameters

- How should candidate hyperparameters be selected for testing? --> Grid search
  - All combinations of selected values of hyperparameters are tested to determine optimal choice.
  - Number of points in grid increases exponentially with number of hyperparameters.
  - Ex- 5 hyperparameters, and test 10 values for each hyperparameter, training procedure needs to be executed  $10^5 = 100000$  times to test its accuracy.
- Grid-based hyperparameter exploration is NOT best choice.

# Hyper-parameters

- Make random sample of hyperparameters uniformly within grid range.
- Logarithms of hyperparameters are sampled uniformly rather than hyperparameters themselves (regularization rate and learning rate).
- Ex: instead of sampling learning rate ( $\alpha$ ) between 0.1 and 0.001, first sample  $\log(\alpha)$  uniformly between  $-1$  and  $-3$ , and then exponentiate it to power of 10.
- Common to search in logarithmic space

# Hyper-parameter Optimization Techniques

- **Babysitting** or ‘Trial and Error’ or Gradient student descent (GSD) method is implemented by 100% manual tuning and widely used.
  - Problems due to large number of hyper-parameters, complex models, time-consuming model evaluations, and non-linear hyper-parameter interactions.
- **Grid search** (GS) most commonly-used methods
  - An exhaustive search or a brute-force method
  - Works by evaluating the Cartesian product of user-specified finite set of values
  - Problem: Inefficiency for high-dimensionality hyper-parameter configuration space
    - Since number of evaluations increases exponentially to number of hyper-parameters growth

# Hyper-parameter Optimization Techniques

- **Random search (RS)**: like GS.
  - Randomly selects a pre-defined number of samples between upper and lower bounds as candidate hyper-parameter values, and then trains these candidates until defined budget is exhausted.
  - With a limited budget, RS can explore **a larger search space than GS**.
  - Problem: unnecessary function evaluations since it does not exploit previously well-performing regions.
- **Gradient-based optimization**: traditional technique.
  - After randomly selecting a data point, it moves towards opposite direction of largest gradient to locate next data point.
    - Local optimum can be reached after convergence.
  - Gradient-based algorithms have a fast convergence speed to reach local optimum.
  - Can use to optimize learning rate in neural networks (NN).

# Hyper-parameter Optimization Techniques Comparisons

The comparison of common HPO algorithms ( $n$  is the number of hyper-parameter values and  $k$  is the number of hyper-parameters).

HPO Method	Strengths	Limitations	Time Complexity
GS	<ul style="list-style-type: none"> <li>• Simple.</li> </ul>	<ul style="list-style-type: none"> <li>• Time-consuming</li> <li>• Only efficient with categorical HPs.</li> </ul>	$O(n^k)$
RS	<ul style="list-style-type: none"> <li>• More efficient than GS</li> <li>• Enable parallelization.</li> </ul>	<ul style="list-style-type: none"> <li>• Not consider previous results</li> <li>• Not efficient with conditional HPs.</li> </ul>	$O(n)$
Gradient-based models	<ul style="list-style-type: none"> <li>• Fast convergence speed for continuous HPs.</li> </ul>	<ul style="list-style-type: none"> <li>• Only support continuous HPs</li> <li>• May only detect local optimums.</li> </ul>	$O(n^k)$
BO-GP	<ul style="list-style-type: none"> <li>• Fast convergence speed for continuous HPs.</li> </ul>	<ul style="list-style-type: none"> <li>• Poor capacity for parallelization</li> <li>• Not efficient with conditional HPs.</li> </ul>	$O(n^3)$
SMAC	<ul style="list-style-type: none"> <li>• Efficient with all types of HPs.</li> </ul>	<ul style="list-style-type: none"> <li>• Poor capacity for parallelization.</li> </ul>	$O(n \log n)$
BO-TPE	<ul style="list-style-type: none"> <li>• Efficient with all types of HPs</li> <li>• Keep conditional dependencies.</li> </ul>	<ul style="list-style-type: none"> <li>• Poor capacity for parallelization.</li> </ul>	$O(n \log n)$
Hyperband	<ul style="list-style-type: none"> <li>• Enable parallelization.</li> </ul>	<ul style="list-style-type: none"> <li>• Not efficient with conditional HPs</li> <li>• Require subsets with small budgets to be representative.</li> </ul>	$O(n \log n)$
BOHB	<ul style="list-style-type: none"> <li>• Efficient with all types of HPs</li> <li>• Enable parallelization.</li> </ul>	<ul style="list-style-type: none"> <li>• Require subsets with small budgets to be representative.</li> </ul>	$O(n \log n)$
GA	<ul style="list-style-type: none"> <li>• Efficient with all types of HPs</li> <li>• Not require good initialization.</li> </ul>	<ul style="list-style-type: none"> <li>• Poor capacity for parallelization.</li> </ul>	$O(n^2)$
PSO	<ul style="list-style-type: none"> <li>• Efficient with all types of HPs</li> <li>• Enable parallelization.</li> </ul>	<ul style="list-style-type: none"> <li>• Require proper initialization.</li> </ul>	$O(n \log n)$

# Hyper-parameter Optimization Frameworks

- Sklearn:
  - GridSearchCV can be implemented to detect optimal hyper-parameters using GS algorithm.
  - RandomizedSearchCV to implement a RS method.
- TensorFlow
  - Trieste: Bayesian optimization toolbox built on TensorFlow  
(<https://github.com/secondmind-labs/trieste>)